

AIM:

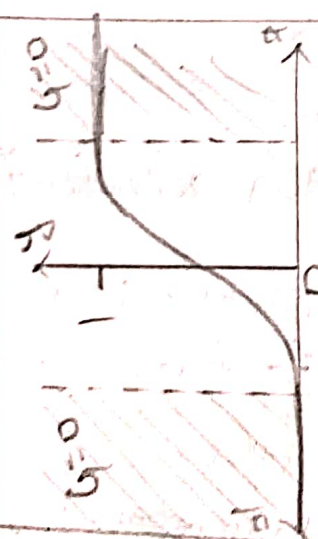
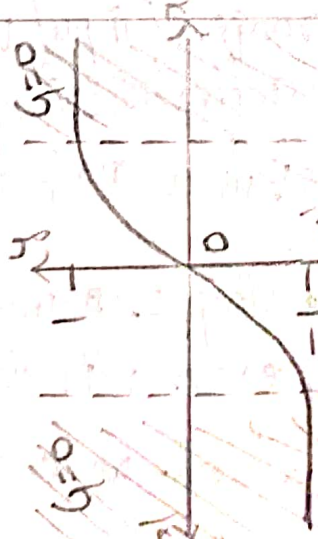
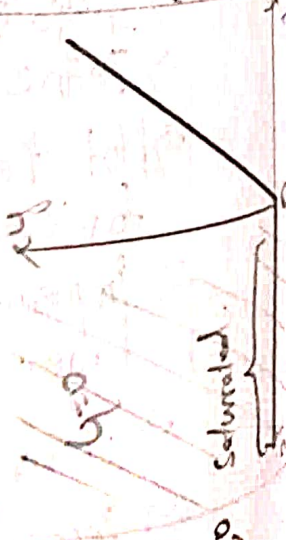
To study various activation functions such as Sigmoid, Tanh, ReLU, Leaky ReLU, and ELU and model performance in neural networks.

OBJECTIVES:

- To understand how activation function introduces non-linearity into neural networks.
- To implement and compare commonly used activation functions.
 - * Sigmoid
 - * Tanh
 - * ReLU (Rectified Linear Unit)
 - * Leaky ReLU
 - * ELU (Exponential Linear Unit)
- To observe their effects on gradient flow and training stability.
- To analyze which activation functions are better suited for different types of problems.

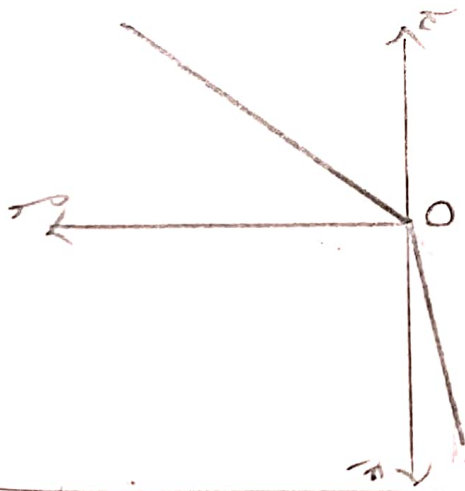
PSEUDOCODE:

- Initialize a range of input values.
- Define each activation function.
- Plot the input values against each activation function to visualize.
- Implement a simple neural network with one layer using each activation function.
- Train the network on sample data.
- Record accuracy, loss, and convergence rate for each function.
- Analyze and compare the behavior based on the observations.

Activation Function	Shape Characteristics	Vanishing Gradient	Computational Cost	Output Range	Shape	Best Use Cases
Sigmoid $\sigma(x) = \frac{1}{1 + e^{-x}}$		Yes	Moderate	(0, 1)	S-shape curve	Binary classification output + layers
Tanh $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$		Yes	Moderate	(-1, 1)	Steeper S-curve	Hidden layers centered data
ReLU $f(x) = \begin{cases} 0, & x \leq 0 \\ x, & x > 0 \end{cases}$		No	Low	[0, ∞)	Linear for positive	" " activation space

Leaky ReLU

$$f(x) = \begin{cases} x & x \geq 0 \\ \alpha x & x < 0 \end{cases}$$



Reduced

low

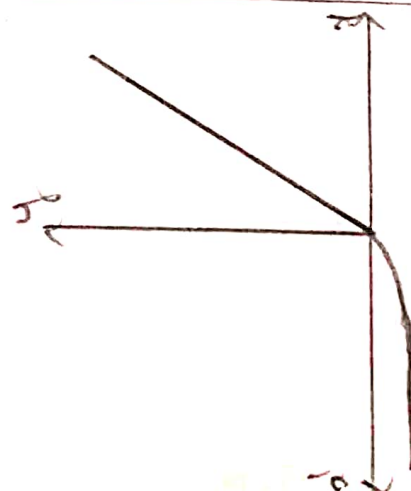
$(-\infty, \infty)$

Allows small slope.

Avoid dying ReLU problem

ELU

$$\max(w_1 x + b_1, w_2 x + b_2, \dots)$$



Reduced

Moderate $(-\infty, \infty)$

Smooth for negative.

Improve learning faster convergence

OBSERVATION:

① Sigmoid

- The training loss decreases very slowly compared to others suffers from slower convergence.
- It starts around 0.92 and ends slightly above 0.6 after 1000 epochs.

② Tanh

- The loss decreases faster than Sigmoid.
- It shows a consistent downward trend. ends at 0.32

③ ReLU

- ReLU starts with a similar initial loss but decreases steadily and converges better than Sigmoid. ends at 0.37

④ Leaky ReLU:

- It has a consistent decline in loss and ends near 0.31, indicating better learning capability than ReLU.

⑤ ELU

- It shows a steady and significant decrease, reaching the lowest loss around 0.31.
- converges faster and better compared.

RESULT:

✓ The experiment was implemented and obtained result successfully.

```
[ ] import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

[ ] # Activation functions and their derivatives
def sigmoid(x):
    return 1 / (1 + np.exp(-x))


[ ] def sigmoid_derivative(x):
    s = sigmoid(x)
    return s * (1 - s)

[ ] def tanh(x):
    return np.tanh(x)

[ ] def tanh_derivative(x):
    return 1 - np.tanh(x)**2

[ ] def relu(x):
    return np.maximum(0, x)

[ ] def relu_derivative(x):
    return (x > 0).astype(float)
```

```
[ ] def tanh_derivative(x):  
    return 1 - np.tanh(x)**2  
  
[ ] def relu(x):  
    return np.maximum(0, x)  
  
[ ] def relu_derivative(x):  
    return (x > 0).astype(float)  
  
[ ] def leaky_relu(x, alpha=0.01):  
    return np.where(x > 0, x, alpha * x)  
  
[ ] def leaky_relu_derivative(x, alpha=0.01):  
    dx = np.ones_like(x)  
    dx[x < 0] = alpha  
    return dx  
  
[ ]  def elu(x, alpha=1.0):  
    return np.where(x >= 0, x, alpha * (np.exp(x) - 1))  
  
[ ] def elu_derivative(x, alpha=1.0):  
    return np.where(x >= 0, 1, elu(x, alpha) + alpha)  
  
[ ] # Simple 1-hidden-layer Neural Network class  
class SimpleNN:
```

lab_5.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Connect

def elu(x, alpha=1.0):
 return np.where(x >= 0, x, alpha * (np.exp(x) - 1))

def elu_derivative(x, alpha=1.0):
 return np.where(x >= 0, 1, elu(x, alpha) + alpha)

Simple 1-hidden-layer Neural Network class

class SimpleNN:
 def __init__(self, input_size, hidden_size, output_size, activation, activation_deriv):
 self.activation = activation
 self.activation_deriv = activation_deriv

 # Initialize weights with small random values
 self.W1 = np.random.randn(input_size, hidden_size) * 0.1
 self.b1 = np.zeros((1, hidden_size))
 self.W2 = np.random.randn(hidden_size, output_size) * 0.1
 self.b2 = np.zeros((1, output_size))

 def forward(self, X):
 self.z1 = X.dot(self.W1) + self.b1
 self.a1 = self.activation(self.z1)
 self.z2 = self.a1.dot(self.W2) + self.b2
 # Output layer - sigmoid for binary classification
 self.a2 = sigmoid(self.z2)
 return self.a2

 def backward(self, X, y, output, lr=0.01):
 m = y.shape[0]

Variables Terminal


```
[ ] ▶  
    return self.a2  
  
def backward(self, X, y, output, lr=0.01):  
    m = y.shape[0]  
  
    # Compute output error  
    delta2 = (output - y) * (output * (1 - output)) # derivative of sigmoid output layer  
  
    # Compute error for hidden layer  
    delta1 = delta2.dot(self.W2.T) * self.activation_deriv(self.z1)  
  
    # Update weights and biases  
    self.W2 -= lr * self.a1.T.dot(delta2) / m  
    self.b2 -= lr * np.sum(delta2, axis=0, keepdims=True) / m  
    self.W1 -= lr * X.T.dot(delta1) / m  
    self.b1 -= lr * np.sum(delta1, axis=0, keepdims=True) / m  
  
def train(self, X, y, epochs=1000, lr=0.01):  
    losses = []  
    for epoch in range(epochs):  
        output = self.forward(X)  
        loss = np.mean(-(y * np.log(output + 1e-9) + (1 - y) * np.log(1 - output + 1e-9)))  
        losses.append(loss)  
        self.backward(X, y, output, lr)  
    return losses  
  
[ ]  
# Generate a toy dataset  
X, y = make_moons(n_samples=500, noise=0.2, random_state=42)  
y = y.reshape(-1, 1)
```



```
[ ] loss = np.mean(-(y * np.log(output + 1e-9) + (1 - y) * np.log(1 - output + 1e-9)))
    losses.append(loss)
    self.backward(X, y, output, lr)
    return losses

[ ] # Generate a toy dataset
X, y = make_moons(n_samples=500, noise=0.2, random_state=42)
y = y.reshape(-1, 1)

[ ] # Scale features
scaler = StandardScaler()
X = scaler.fit_transform(X)

[ ] # Split train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

[ ] # Dictionary of activation functions and derivatives to test
activations = {
    'Sigmoid': (sigmoid, sigmoid_derivative),
    'Tanh': (tanh, tanh_derivative),
    'ReLU': (relu, relu_derivative),
    'Leaky ReLU': (leaky_relu, leaky_relu_derivative),
    'ELU': (elu, elu_derivative)
}

[ ] # Train networks with different activations and store losses
losses_dict = {}
```

Indian S | ISRO Lc | GATE 20 | GATE 20 | GATE | GATE20 | Inbox (1) | Inbox (1) | UROP | 21AIC30 | Non-re | drive - S | Colab N | lab x | lab_6 - | Lab 5 a |

https://colab.research.google.com/drive/1CX6pArhhDo7azAKmE98BbLXS140Sxa0e?authuser=2

lab_5.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Connect

[]

Sigmoid : (sigmoid, sigmoid_derivative),
'Tanh': (tanh, tanh_derivative),
'ReLU': (relu, relu_derivative),
'Leaky ReLU': (leaky_relu, leaky_relu_derivative),
'ELU': (elu, elu_derivative)
}

[]

Train networks with different activations and store losses
losses_dict = {}

[]

for name, (act, act_deriv) in activations.items():
 print(f"Training with {name} activation...")
 nn = SimpleNN(input_size=2, hidden_size=10, output_size=1, activation=act, activation_deriv=act_deriv)
 losses = nn.train(X_train, y_train, epochs=1000, lr=0.05)
 losses_dict[name] = losses

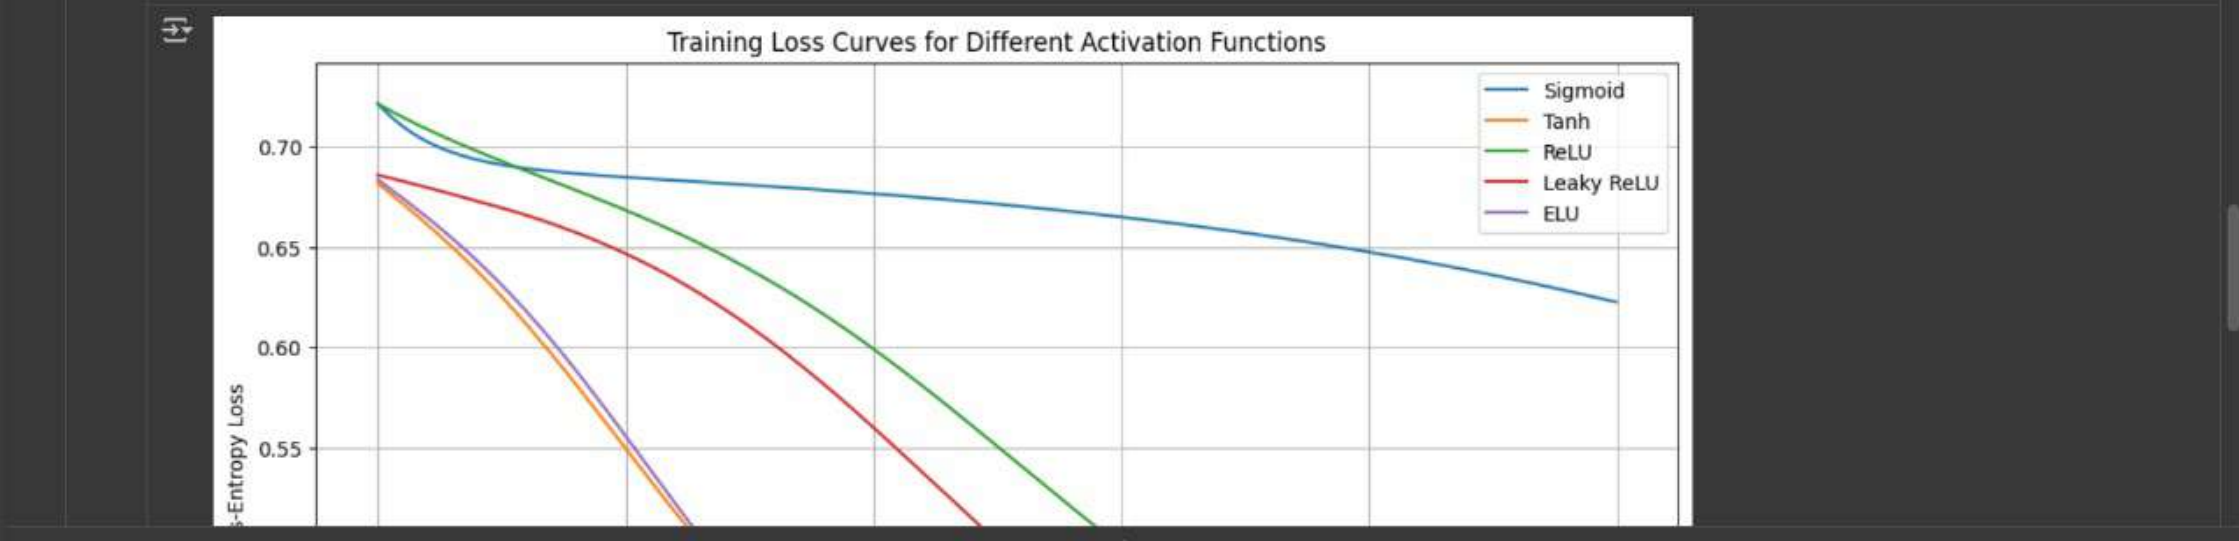
Training with Sigmoid activation...
Training with Tanh activation...
Training with ReLU activation...
Training with Leaky ReLU activation...
Training with ELU activation...

[]

Plot the loss curves for comparison
plt.figure(figsize=(12, 8))
for name, losses in losses_dict.items():
 plt.plot(losses, label=name)
plt.title("Training Loss Curves for Different Activation Functions")
plt.xlabel("Epochs")
plt.ylabel("Binary Cross-Entropy Loss")
plt.legend()

Variables Terminal

```
# Plot the loss curves for comparison
plt.figure(figsize=(12, 8))
for name, losses in losses_dict.items():
    plt.plot(losses, label=name)
plt.title("Training Loss Curves for Different Activation Functions")
plt.xlabel("Epochs")
plt.ylabel("Binary Cross-Entropy Loss")
plt.legend()
plt.grid(True)
plt.show()
```



Training Loss Curves for Different Activation Functions

