




Loss vs Epochs



pc (イリコ (い。 ,



Exp-7 BUILD A CNN MODEL TO CLASSIFY CAT & DOG IMAGES

AIM


To design and implement a Convolution Neural Network (CNN) model to classify images of cats and dogs using a binary classification approach.

OBJECTIVES

- Load and preprocess the dataset of cat and dog images.
- Build a CNN model for binary image classification.
- Train and validate the model using appropriate evaluation metrics.
- Plot training history (loss vs accuracy).
- Evaluate model performance using confusion matrix, classification report, and ROC-AUC curve.

PSEUDO CODE

- Load and preprocess image dataset:
 - Resize image to uniform size
 - Normalize pixel values
 - Split into training and validation sets.
- Building CNN model:
 - Add CONV2D, MAX Pooling 2D, Dropout and Dense layers.
 - Use sigmoid activation for binary classification
- Compile the model:
 - Use 'binary-crossentropy' as loss
 - Use 'adam' optimizer
 - Track 'accuracy' metric



→ Train the model:

- Use training and validation data

- Set epochs and batch size

→ Evaluate the model:

- Plot accuracy and loss curves

- Plot ROC-AUC curve.

- Print classification report and confusion matrix

OBSERVATION

	Precision	recall	f1-score	support
Cat	0.84	0.80	0.82	2527
Dog	0.81	0.84	0.83	2465
accuracy			0.82	4992
macroavg	0.82	0.82	0.82	4992
weighted avg	0.82	0.82	0.82	4992

Confusion Matrix

Actual	Cat	Dog
Cat	2191	331
Dog	516	1956
Predicted	Cat	Dog

RESULT:

The CNN model is implemented and we have got the accuracy of 82% successfully.

https://colab.research.google.com/drive/17Jv3PPIhdIJ1AigFBgoN8KIZ5Ejv_dS?authuser=2

LAB_7.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect T4

```
import kagglehub
# Download latest version
path = kagglehub.dataset_download("karakaggle/kaggle-cat-vs-dog-dataset")

print("Path to dataset files:", path)

Downloading from https://www.kaggle.com/api/v1/datasets/download/karakaggle/kaggle-cat-vs-dog-dataset?dataset_version_number=1...
100%|██████████| 787M/787M [00:08:00.00, 93.8MB/s]Extracting files...
Path to dataset files: /root/.cache/kagglehub/datasets/karakaggle/kaggle-cat-vs-dog-dataset/versions/1
```

```
import os
os.chdir('/root/.cache/kagglehub/datasets/karakaggle/kaggle-cat-vs-dog-dataset/versions/1')

os.listdir()
['kagglecatsanddogs_3367a']

os.chdir('kagglecatsanddogs_3367a')

os.listdir()
['readme[1].txt', 'MSR-LA - 3467.docx', 'PetImages']
```

```
nc_chdir("PetImages")
```

Variables Terminal

Academia - Academic Web Servic... LAB_7.ipynb - Colab DLT_LAB/OBSERVATION at main

https://colab.research.google.com/drive/17Jv3PPIhdIJ1AigFBgoN8KIZ5Ejv_dS?authuser=2

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect T4

```
['readme[1].txt', 'MSR-LA - 3467.docx', 'PetImages']
```

```
os.chdir("PetImages")
```

```
os.getcwd()
```

```
'/root/.cache/kagglehub/datasets/karakaggle/kaggle-cat-vs-dog-dataset/versions/1/kagglecatsanddogs_3367a/PetImages'
```

```
import os
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
from torchvision import datasets, transforms, models
```

```
# Set device (GPU if available)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

Using device: cuda
```

```
# Parameters
```

Variables Terminal

Academia - Academic Web Servic... LAB_7.ipynb - Colab DLT_LAB/OBSERVATION at main

https://colab.research.google.com/drive/17Jv3PPIhdIJ1AigFBgoN8KIZ5Ejv_dS?authuser=2

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect T4

```
Using device: cuda
```

```
# Parameters
DATA_DIR = '/root/.cache/kagglehub/datasets/karakaggle/kaggle-cat-vs-dog-dataset/versions/1/kagglecatsanddogs_3367a/PetImages'
BATCH_SIZE = 32
IMG_SIZE = 150
EPOCHS = 10
LR = 1e-4
# Data Transforms
train_transform = transforms.Compose([
    transforms.Resize(IMG_SIZE, IMG_SIZE),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(20),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

val_transform = transforms.Compose([
    transforms.Resize(IMG_SIZE, IMG_SIZE),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

# Load Dataset
full_dataset = datasets.ImageFolder(DATA_DIR, transform = train_transform)

# Train-Val Split (80%-20%)
train_size = int(0.8 * len(full_dataset))
```

Variables Terminal

Academia - Academic Web Service LAB_7.ipynb - Colab DLT_LAB/OBSERVATION at main...

```
# Load Dataset
full_dataset = datasets.ImageFolder(DATA_DIR, transform = train_transform)

# Train-Val Split (80%-20%)
train_size = int(0.8 * len(full_dataset))
val_size = len(full_dataset) - train_size
train_dataset, val_dataset = random_split(full_dataset, [train_size, val_size])

# Change val dataset transform
val_dataset.dataset.transform = val_transform

# DataLoaders
train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)
```

Variables Terminal

Academia - Academic Web Service LAB_7.ipynb - Colab DLT_LAB/OBSERVATION at main...

```
# DataLoaders
train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)

# Simple CNN Model
class SimpleCNN(nn.Module):
    def __init__(self):
        super(SimpleCNN, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1), # (B,32,150,150)
            nn.ReLU(),
            nn.MaxPool2d(2), # (B,32,75,75)
```

Variables Terminal

Academia - Academic Web Service LAB_7.ipynb - Colab DLT_LAB/OBSERVATION at main...

```
x = self.features(x)
x = self.classifier(x)
return x

model = SimpleCNN().to(device)
print(model)

SimpleCNN:
(features): Sequential(
(0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(1): ReLU()
(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(4): ReLU()
(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(6): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(7): ReLU()
(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Flatten(start_dim=1, end_dim=-1)
(1): Dropout(p=0.5, inplace=False)
(2): Linear(in_features=41472, out_features=128, bias=True)
(3): ReLU()
(4): Linear(in_features=128, out_features=1, bias=True)
(5): Sigmoid()
)
```

Loss and optimizer
criterion = nn.BCELoss()

Academia - Academic Web Servi... LAB_7.ipynb - Colab DLT_LAB/OBSERVATION at main...

loss and optimizer
criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=LR)

Training and validation loop
train_losses, val_losses = [], []
train_accs, val_accs = [], []

for epoch in range(EPOCHS):
model.train()
running_loss, running_corrects = 0, 0

for inputs, labels in train_loader:
inputs = inputs.to(device)
labels = labels.unsqueeze(1).float().to(device) # For BCELoss

optimizer.zero_grad()
outputs = model(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

running_loss += loss.item() * inputs.size(0)
preds = (outputs > 0.5).float()
running_corrects += torch.sum(preds == labels)

epoch_loss = running_loss / train_size
epoch_acc = running_corrects.double() / train_size
train_losses.append(epoch_loss)
train_accs.append(epoch_acc.item())

Validation
model.eval()
val_running_loss, val_running_corrects = 0, 0
all_preds = []
all_labels = []

with torch.no_grad():
for inputs, labels in val_loader:
inputs = inputs.to(device)
labels = labels.unsqueeze(1).float().to(device)
outputs = model(inputs)
loss = criterion(outputs, labels)

val_running_loss += loss.item() * inputs.size(0)

val_running_corrects += torch.sum(preds == labels)
all_preds.extend(outputs.cpu().numpy())
all_labels.extend(labels.cpu().numpy())


val_loss = val_running_loss / val_size
val_acc = val_running_corrects.double() / val_size
val_losses.append(val_loss)
val_accs.append(val_acc.item())

print(f"Epoch {epoch+1}/{EPOCHS}..")
f"Train loss: {epoch_loss:.4f}.. Train acc: {epoch_acc:.4f}.."
f"Val loss: {val_loss:.4f}.. Val acc: {val_acc:.4f}")

```
/usr/local/lib/python3.12/dist-packages/PIL/TiffImagePlugin.py:950: UserWarning: Truncated File Read
warnings.warn(str(msg))
```

Epoch 1/10.. Train loss: 0.6069.. Train acc: 0.6664.. Val loss: 0.5304.. Val acc: 0.7396
Epoch 2/10.. Train loss: 0.5138.. Train acc: 0.7474.. Val loss: 0.4719.. Val acc: 0.7758
Epoch 3/10.. Train loss: 0.4651.. Train acc: 0.7810.. Val loss: 0.4633.. Val acc: 0.7736
Epoch 4/10.. Train loss: 0.4291.. Train acc: 0.8036.. Val loss: 0.4268.. Val acc: 0.8049
Epoch 5/10.. Train loss: 0.4008.. Train acc: 0.8200.. Val loss: 0.3986.. Val acc: 0.8195
Epoch 6/10.. Train loss: 0.3748.. Train acc: 0.8314.. Val loss: 0.4084.. Val acc: 0.8185
Epoch 7/10.. Train loss: 0.3520.. Train acc: 0.8471.. Val loss: 0.4005.. Val acc: 0.8103
Epoch 8/10.. Train loss: 0.3288.. Train acc: 0.8578.. Val loss: 0.3882.. Val acc: 0.8233
Epoch 9/10.. Train loss: 0.3098.. Train acc: 0.8640.. Val loss: 0.3769.. Val acc: 0.8303
Epoch 10/10.. Train loss: 0.2895.. Train acc: 0.8742.. Val loss: 0.3741.. Val acc: 0.8303

Plot loss and accuracy curves
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)




```
# Evaluate with sklearn metrics
all_preds = np.array(all_preds).flatten()
all_preds_binary = (all_preds > 0.5).astype(int)
all_labels = np.array(all_labels).flatten()

print("Classification Report:\n", classification_report(all_labels, all_preds_binary, target_names=full_dataset.classes))

Classification Report:
precision    recall   f1-score   support
Cat          0.81     0.87     0.84      2522
Dog          0.86     0.79     0.82      2478

accuracy           0.83     0.83     0.83      4992
macro avg       0.83     0.83     0.83      4992
weighted avg    0.83     0.83     0.83      4992

# Confusion matrix plot
cm = confusion_matrix(all_labels, all_preds_binary)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=full_dataset.classes, yticklabels=full_dataset.classes)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



Academia - Academic Web Servic x LAB_7.ipynb - Colab x DLT_LAB/OBSERVATION at main · +

https://colab.research.google.com/drive/17JV3PPihdlI1AigFBgoN8KIZ5Ejv_dS?authuser=2

LAB_7.ipynb Share K

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

```
[1]: # ROC-AUC Curve
fpr, tpr, _ = roc_curve(all_labels, all_preds)
roc_auc = roc_auc_score(all_labels, all_preds)

plt.figure(figsize=(6,6))
plt.plot(fpr, tpr, label=f'ROC AUC = {roc_auc:.2f}')
plt.plot([0,1],[0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC-AUC Curve')
plt.legend(loc='lower right')
plt.grid()
plt.show()
```

ROC-AUC Curve

Positive Rate

