

# ResNet - 18

Input

$3 \times 3 \text{ conv. } 64$

$3 \times 3 \text{ conv. } 128, 128$

$3 \times 3 \text{ conv. } 256, 256$

$3 \times 3 \text{ conv. } 512, 512$

$3 \times 3 \text{ conv. } 512$

$3 \times 3 \text{ conv. } 512$

$3 \times 3 \text{ conv. } 512$

Avg pool  $1 \times 1$

↓  
FC



Exp: 13

## UNDERSTANDING THE ARCHITECTURE OF PRE-TRAINED MODEL

AIM

To explore and analyse the architecture of pre-trained convolutional neural network ResNet-18 visualize its layers, parameters, feature maps, and understand its behavior on sample images.



OBJECTIVE:

- \* Load a pre-trained model from PyTorch's torchvision.models
- \* Examine the layer structure and number of parameters
- \* Visualize feature maps from intermediate convolutional layers.
- \* Understand how low-level and high-level features are extracted.
- \* Compute and visualize accuracy and loss metrics on small dataset.

PSEUDO CODE:

- ↳ Import required libraries
- ↳ Load a pre-trained model (ResNet-18) and set to eval mode.
- ↳ Print the model summary and total parameters.
- ↳ Forward pass image through model
  - Extract output predictions
  - Extract intermediate feature maps
- ↳ Visualize feature maps as grids
- ↳ Optionally load a small dataset (CIFAR-10 subset).
- ↳ Plot graphs for loss and accuracy.

## Cross Entropy loss per Batch



Top-1 accuracy is very low, fluctuating slightly around 0.000. This indicates that the model is not learning effectively or that the data is too noisy for the current architecture.

OBSERVATION :

Total parameters: 11689512

Trainable parameter: 11689512

\* Only the final layer is trained; convolutional layers remain feature extractors.

\* Can unfreeze layers later for fine-tuning to further improve performance.

\* Using transfer learning is especially useful when the dataset is small or medium-sized.

RESULT:

Successfully implemented <sup>new</sup> ResNet-18 model

Rajat

```
import torch
from torchvision import models, datasets, transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import matplotlib.pyplot as plt
import numpy as np
```

Device

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

Load pre-trained model

```
model = models.resnet18(pretrained=True).to(device)
model.eval()
```

ResNet(  
 (conv1): Conv2d(3, 64, kernel\_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)  
 (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)  
 (relu): ReLU(inplace=True)  
 (maxpool): MaxPool2d(kernel\_size=3, stride=2, padding=1, dilation=1, ceil\_mode=False)  
 (layer1): Sequential(  
 (0): BasicBlock(  
 (conv1): Conv2d(64, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
 (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)  
 (relu): ReLU(inplace=True)  
 )  
 (1): BasicBlock(  
 (conv1): Conv2d(64, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
 (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)  
 (relu): ReLU(inplace=True)  
 )  
 )  
 (layer2): Sequential(  
 (0): BasicBlock(  
 (conv1): Conv2d(64, 128, kernel\_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)  
 (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)  
 (relu): ReLU(inplace=True)  
 (conv2): Conv2d(128, 128, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
 (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)  
 )  
 (1): BasicBlock(  
 (conv1): Conv2d(128, 128, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
 (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)  
 (relu): ReLU(inplace=True)  
 (conv2): Conv2d(128, 128, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
 (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track\_running\_stats=True)  
 )  
 )  
)

```
(bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(downsample): Sequential(
  (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(1): BasicBlock(
  (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
)
```

Academia - Academic Web Service LAB\_13 - Colab KAVINRAJ06/DLT\_LAB

[https://colab.research.google.com/drive/1ZnGM0t8vdPPFRLsshgScvRaD\\_HGmdk?authuser=2](https://colab.research.google.com/drive/1ZnGM0t8vdPPFRLsshgScvRaD_HGmdk?authuser=2)

**LAB\_13**

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

```
(bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (downsample): Sequential(
    (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc): Linear(in_features=512, out_features=1000, bias=True)
)
```

Prepare dataset (subset of CIFAR-10 for demonstration)

```
[ ] transform = transforms.Compose([
  transforms.Resize((224,224)),
  transforms.ToTensor(),
  transforms.Normalize([0.485,0.456,0.406], [0.229,0.224,0.225])
])

test_dataset = datasets.CIFAR10(root='./keras/datasets', train=False, download=True, transform=transform)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
```

Variables Terminal

Academia - Academic Web Service LAB\_13 - Colab KAVINRAJ06/DLT\_LAB

[https://colab.research.google.com/drive/1ZnGM0t8vdPPFRLsshgScvRaD\\_HGmdk?authuser=2](https://colab.research.google.com/drive/1ZnGM0t8vdPPFRLsshgScvRaD_HGmdk?authuser=2)

**LAB\_13**

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

Prepare dataset (subset of CIFAR-10 for demonstration)

```
[ ] transform = transforms.Compose([
  transforms.Resize((224,224)),
  transforms.ToTensor(),
  transforms.Normalize([0.485,0.456,0.406], [0.229,0.224,0.225])
])

test_dataset = datasets.CIFAR10(root='./keras/datasets', train=False, download=True, transform=transform)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
```

Define loss function

```
[ ] criterion = nn.CrossEntropyLoss()
```

Evaluate model on dataset

```
[ ] all_losses = []
all_accuracies = []

with torch.no_grad():
  for imgs, labels in test_loader:
    imgs, labels = imgs.to(device), labels.to(device)

  outputs = model(imgs) # [batch, 1000]
```

Variables Terminal

Academia - Academic Web Service LAB\_13 - Colab KAVINRAJ06/DLT\_LAB

```

all_losses = []
all_accuracies = []

with torch.no_grad():
    for imgs, labels in test_loader:
        imgs, labels = imgs.to(device), labels.to(device)

        outputs = model(imgs) # [batch, 1000]

        # Compute loss (for demonstration; note labels are CIFAR-10, model trained on ImageNet)
        # So loss values are arbitrary, but can be visualized
        labels_expanded = torch.randint(0, 1000, labels.shape).to(device) # simulate compatible labels
        loss = criterion(outputs, labels_expanded)
        all_losses.append(loss.item())

        # Compute accuracy (simulate top-1 accuracy for visualization)
        preds = outputs.argmax(dim=1)
        acc = (preds == labels_expanded).float().mean().item()
        all_accuracies.append(acc)

```

Plot Loss and Accuracy

```

total_params = sum(p.numel() for p in model.parameters())
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"Total parameters: {total_params}")
print(f"Trainable parameters: {trainable_params}")

Total parameters: 11689512
Trainable parameters: 11689512

```

Variables Terminal

Academia - Academic Web Service LAB\_13 - Colab KAVINRAJ06/DLT\_LAB

```

total_params = sum(p.numel() for p in model.parameters())
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"Total parameters: {total_params}")
print(f"Trainable parameters: {trainable_params}")

Total parameters: 11689512
Trainable parameters: 11689512

```

```

plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plt.plot(all_losses, label="Loss", color='tab:red')
plt.title("Cross-Entropy Loss per Batch")
plt.xlabel("Batch")
plt.ylabel("Loss")
plt.grid(True)
plt.legend()

plt.subplot(1,2,2)
plt.plot(all_accuracies, label="Accuracy", color='tab:blue')
plt.title("Top-1 Accuracy per Batch")
plt.xlabel("Batch")
plt.ylabel("Accuracy")
plt.grid(True)
plt.legend()

```

Variables Terminal

