

IMPLEMENT A YOLO MODEL TO DETECT OBJECT

AIM

To develop a robust traffic sign detection system using YOLOv8 that can accurately identify and locate multiple traffic signs in images/videos.

OBJECTIVES:

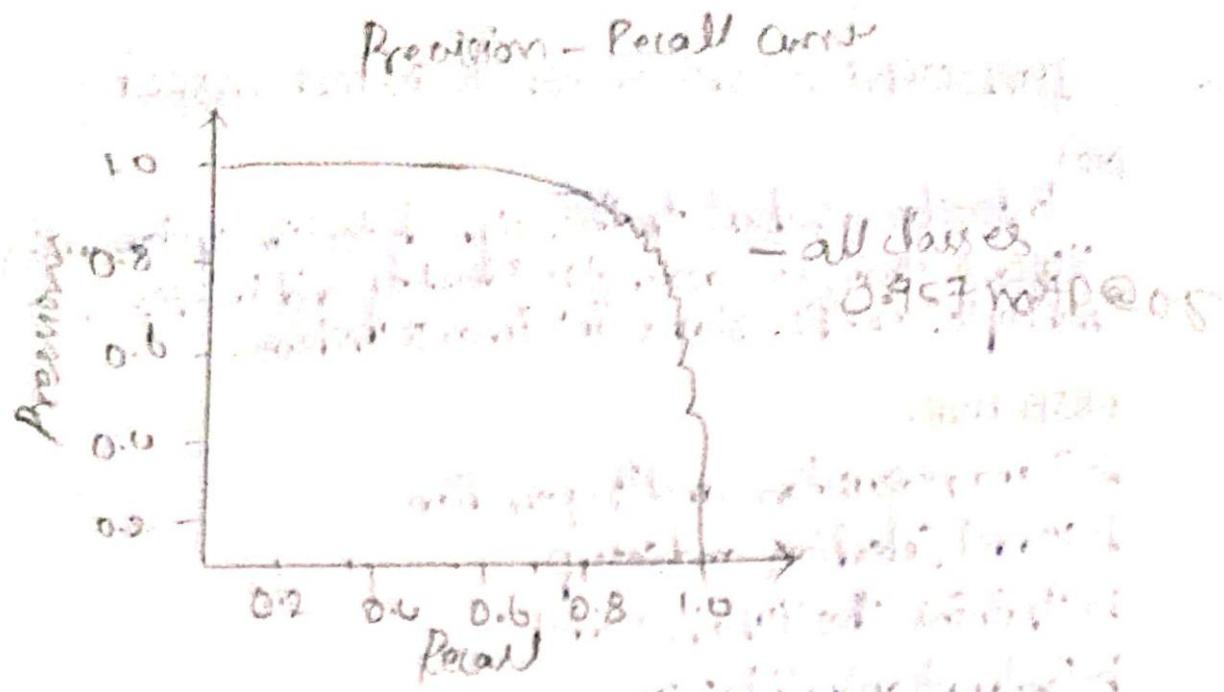
- ↳ Data Acquisition and Preparation
- ↳ Model Selection and setup
- ↳ Training the YOLOv8 Model
- ↳ Evaluation Validation.
- ↳ Prediction and Inference
- ↳ Model Export and Deployment.

PSEUDOCODE:

- ↳ Import necessary libraries
- ↳ Download and prepare the dataset.
- ↳ Initialize YOLOv8 model
- ↳ Train YOLOv8 model
- ↳ Evaluate model performance
- ↳ Perform inference on test images
 - for each image in testset:
- ↳ Perform inference on videos.
- ↳ Export final model.
- ↳ Deployment (optional)

OBSERVATIONS:

Epoch	box-loss	cls-loss	IoU	size
1/30	0.9003	3.816	1.231	640
2/30	0.7314	2.68	106	640
3/30	0.7322	1.916	121	640
4/30	0.7114	1.914	121	640
5/30	0.6686	1.391	114	640



precision-recall graph
should be compared to recall of
random classifier (diagonal line)
from recall about 0.5
down to about 0.15
is good, if recall drops
below 0.5 it is difficult to
decide if predicted correctly
because it depends on
thresholds.

	Random	Naive Bayes	Logistic	SVM	Decision Tree
Recall	0.00	0.00	0.00	0.00	0.00
Precision	0.00	0.00	0.00	0.00	0.00
F1 Score	0.00	0.00	0.00	0.00	0.00
AUC	0.50	0.50	0.50	0.50	0.50

29/30	0.4945	0.3838	0.9014	65	640
30/30	0.4902	0.3806	0.9008	188	640

Precision : 0.9493981129777693

metrues | recall

metrues | mAP50(B) : 0.9060281630136168

metrues | mAP50-95(B) : 0.9571950392546826

metrues | mAP50-95(B) : 0.8296656258388937.

~~Result~~

(E) Successfully Implemented a YOLOv8 model to
detect object.

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search Select Kernel

LAB_15 (1).ipynb

C:\Users\ASUS\Downloads > LAB_15 (1).ipynb > M4 Install And Import Essential Libraries

Generate + Code + Markdown | Run All Clear All Outputs | Outline ...

Install And Import Essential Libraries

```
# Install Essential Libraries
!pip install ultralytics
```

[1]

Requirement already satisfied: ultralytics in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: numpy>=1.23.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: matplotlib>=3.3.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: opencv-python>=4.6.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: pillow>=7.1.2 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: pyyaml>=5.3.1 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: requests>=2.23.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: scipy>=1.4.1 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: torch>=1.8.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: torchvision>=0.9.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: psutil in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages (from
Requirement already satisfied: polaris in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: ultralytics-thop>=2.0.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: contourpy>=1.0.1 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: cython>=0.10 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: fonttools>=4.22.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: kiwicrawler>=1.3.1 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: packaging>=20.0 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: python-dateutil>=2.7 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: six>=1.5 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages (fr
Requirement already satisfied: charset-normalizer<4,>>2 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: idna<4,>>2.5 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: urllib3<3,>>1.21.1 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
Requirement already satisfied: certifi>=2017.4.17 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages
...
Requirement already satisfied: sympy==1.13.1 in c:\users\asus\appdata\local\packages\pythonsoftwarefoundation.python.3.10_qbz5n2kfra8p0\localcache\local_packages\python310\site-packages

Cell 1 of 20

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search Select Kernel

LAB_15 (1).ipynb

C:\Users\ASUS\Downloads > LAB_15 (1).ipynb > M4 Install And Import Essential Libraries

Generate + Code + Markdown | Run All Clear All Outputs | Outline ...

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output settings...

[notice] A new release of pip is available: 25.1.1 -> 25.3.1
[notice] To update, run: C:\Users\ASUS\AppData\Local\Microsoft\WindowsApps\PythonSoftwareFoundation.Python.3.10_qbz5n2kfra8p0\python.exe -m pip install --upgrade pip

```
# Import Essential Libraries
import os
import random
import pandas as pd
from PIL import Image
import cv2
from ultralytics import YOLO
from IPython.display import Video
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style='darkgrid')
import pathlib
import glob
from tqdm.notebook import trange, tqdm
import warnings
warnings.filterwarnings('ignore')
```

[2]

```
# Configure the visual appearance of Seaborn plots
sns.set(rc={'axes.facecolor': '#eae8fa'}, style='darkgrid')
```

[3]

Cell 1 of 20

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search ⏴ ⏵

C:\Users\ASUS\Downloads > LAB_15 (1).ipynb > Install And Import Essential Libraries

Generate + Code + Markdown | Run All | Clear All Outputs | Outline ... Select Kernel Python

```
# Configure the visual appearance of Seaborn plots
sns.set(rc={'axes.facecolor': '#eaeafa'}, style='darkgrid')
```

[3] Dataset

```
import kagglehub

# Download latest version
path = kagglehub.dataset_download("pkdarabi/cardetection")

print("Path to dataset files:", path)
```

[4] Path to dataset files: C:\Users\ASUS\.cache\kagglehub\datasets\pkdarabi\cardetection\versions\5

```
Image_dir = r'D:\car\train\images'

num_samples = 9
image_files = os.listdir(Image_dir)

# Randomly select num_samples images
rand_images = random.sample(image_files, num_samples)

fig, axes = plt.subplots(3, 3, figsize=(11, 11))

for i in range(num_samples):
    image = rand_images[i]
    ax = axes[i // 3, i % 3]
    ax.imshow(plt.imread(os.path.join(Image_dir, image)))
```

[5] Cell 1 of 20

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search ⏴ ⏵

C:\Users\ASUS\Downloads > LAB_15 (1).ipynb > Install And Import Essential Libraries

Generate + Code + Markdown | Run All | Clear All Outputs | Outline ... Select Kernel Python

```
Image_dir = r'D:\car\train\images'

num_samples = 9
image_files = os.listdir(Image_dir)

# Randomly select num_samples images
rand_images = random.sample(image_files, num_samples)

fig, axes = plt.subplots(3, 3, figsize=(11, 11))

for i in range(num_samples):
    image = rand_images[i]
    ax = axes[i // 3, i % 3]
    ax.imshow(plt.imread(os.path.join(Image_dir, image)))
    ax.set_title(f'Image {i+1}')
    ax.axis('off')

plt.tight_layout()
plt.show()
```

[6] Cell 1 of 20

LAB_15 (1).ipynb

C:\Users\ASUS\Downloads > LAB_15 (1).ipynb > **M4 Install And Import Essential Libraries**

Generate + Code + Markdown | Run All Clear All Outputs | Outline ... Select Kernel

Image 4: A speed limit sign showing 60 km/h. The image is slightly blurred.

Image 5: A speed limit sign showing 20 km/h. The image is heavily distorted with colorful artifacts.

Image 6: A speed limit sign showing 80 km/h. The image is clear and shows a road with other vehicles in the background.

Image 4: A speed limit sign showing 60 km/h. The image is heavily distorted with colorful artifacts.

Image 5: A speed limit sign showing 20 km/h. The image is clear and shows a road with other vehicles in the background.

Image 6: A speed limit sign showing 80 km/h. The image is clear and shows a road with other vehicles in the background.

Image 7: A speed limit sign showing 100 km/h. The image is clear and shows a road along a cliffside.

Image 8: A speed limit sign showing 30 km/h. The image is heavily distorted with colorful artifacts.

Image 9: A speed limit sign showing 40 km/h. The image is clear and shows a road with greenery in the background.

```
# Get the size of the image
image = cv2.imread("D:\car\train\images\000000.jpg")
h, w, c = image.shape
print("The image has dimensions {w}x{h} and {c} channels.")

[Python]
[Python]
... The image has dimensions 416x416 and 3 channels.

Pre-trained YOLOv8 For Detect Traffic Signs
```

Cell 1 of 20

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search Select Kernel Python

```

LAB_15 (1).ipynb x
C:\Users\ASUS\Downloads > LAB_15 (1).ipynb > M+ Install And Import Essential Libraries
Generate + Code + Markdown | Run All Clear All Outputs | Outline ...
# Get the size of the image
image = cv2.imread("D:\car\train\images\000000.jpg")
h, w, c = image.shape
print(f"The image has dimensions {w}x{h} and {c} channels.")

[8] ... The image has dimensions 416x416 and 3 channels.

Pre-trained YOLOv8 For Detect Traffic Signs

# Use a pretrained YOLOv8n model
model = YOLO("yolov8n.pt")

# use the model to detect object
image = r"D:\car\train\images\000000.jpg"
result_predict = model.predict(source = image, imgsz=(640))

# show results
plot = result_predict[0].plot()
plot = cv2.cvtColor(plot, cv2.COLOR_BGR2RGB)
display(Image.fromarray(plot))

[9] ... Downloading https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8n.pt to 'yolov8n.pt': 100% ━━━━━━━━ 6.2MB 56.2KB/s 1:54 1:54<0.2s2:06
image 1/1 D:\car\train\images\000000.jpg_rf_b11f308f16626f9f795a148029c46d10.jpg: 640x640 (no detections), 18.7ms
Speed: 7.8ms preprocess, 18.7ms inference, 2.5ms postprocess per image at shape (1, 3, 640)
...

```



Cell 1 of 20

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search Select Kernel Python

```

LAB_15 (1).ipynb x
C:\Users\ASUS\Downloads > LAB_15 (1).ipynb > M+ Install And Import Essential Libraries
Generate + Code + Markdown | Run All Clear All Outputs | Outline ...
... Downloading https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8n.pt to 'yolov8n.pt': 100% ━━━━━━━━ 6.2MB 56.2KB/s 1:54 1:54<0.2s2:06
image 1/1 D:\car\train\images\000000.jpg_rf_b11f308f16626f9f795a148029c46d10.jpg: 640x640 (no detections), 18.7ms
Speed: 7.8ms preprocess, 18.7ms inference, 2.5ms postprocess per image at shape (1, 3, 640)
...

```



Cell 1 of 20

```

[12] # Build from YAML and transfer weights
from ultralytics import YOLO

```

LAB_15 (1).ipynb

C:\> Users > ASUS > Downloads > LAB_15 (1).ipynb > **Run All**

```
# Build from YAML and transfer weights
from ultralytics import YOLO

Final_model = YOLO('yolov8n.pt')

# Training The Final Model
Result_Final_model = Final_model.train(data=r"D:\car\data.yaml", epochs = 10, batch = -1, optimizer = 'auto')

[12] ... New https://pypi.org/project/ultralytics/8.3.221 available Update with 'pip install -U ultralytics'
Ultralytics 8.3.221 Python 3.10.11 torch-2.5.1cu121 CUDA:0 (NVIDIA GeForce RTX 3090 Laptop GPU, 4096MiB)
engine/trainer: agnostic_nms=False, amp=True, augment=False, auto_augment=randaugment, batch=-1, bgr=0.0, box=7.5, cache=False, cfg=None, classes=None, close_mosaic=10, cls=0.5, compil
Overriding model.yaml nc=80 with nc=15

from params module
arguments
0   -1 1    464 ultralytics.nn.modules.conv.Conv      [3, 16, 3, 2]
1   -1 1    4672 ultralytics.nn.modules.conv.Conv     [16, 32, 3, 2]
2   -1 1    7360 ultralytics.nn.modules.block.C2f     [32, 32, 1, True]
3   -1 1    18560 ultralytics.nn.modules.conv.Conv     [32, 64, 3, 2]
4   -1 2    49664 ultralytics.nn.modules.block.C2f     [64, 64, 2, True]
5   -1 1    73984 ultralytics.nn.modules.conv.Conv     [64, 128, 3, 2]
6   -1 2    197632 ultralytics.nn.modules.block.C2f     [128, 128, 2, True]
7   -1 1    295424 ultralytics.nn.modules.conv.Conv     [128, 256, 3, 2]
8   -1 2    460288 ultralytics.nn.modules.block.C2f     [256, 256, 1, True]
9   -1 1    164088 ultralytics.nn.modules.block.SPPF    [256, 256, 5]
10  -1 1    0 torch.nn.modules.upsampling.Upsample      [None, 2, 'nearest']
11  [-1, 6] 1    0 ultralytics.nn.modules.conv.Concat    [1]
12  -1 1    148224 ultralytics.nn.modules.block.C2f     [384, 128, 1]
13  -1 1    0 torch.nn.modules.upsampling.Upsample      [None, 2, 'nearest']
14  [-1, 4] 1    0 ultralytics.nn.modules.conv.Concat    [1]
15  -1 1    37248 ultralytics.nn.modules.block.C2f      [192, 64, 1]
16  -1 1    36992 ultralytics.nn.modules.conv.Conv     [64, 64, 3, 2]
17  [-1, 12] 1    0 ultralytics.nn.modules.concat.Concat [1]
18  -1 1    123648 ultralytics.nn.modules.block.C2f     [192, 128, 1]

[14] ... Speed Limit 90      38      38      0.275      0.632      0.358      0.307
[15] Stop      81      81      0.945      0.84      0.924      0.826
Speed: 0.4ms preprocess, 5.0ms inference, 0.0ms loss, 1.2ms postprocess per image
Results saved to C:\Users\ASUS\Downloads\runs\detect\train3
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```

Validation Step

```
import os
import cv2
import matplotlib.pyplot as plt

def display_images(post_training_files_path, image_files):
    for image_file in image_files:
        image_path = os.path.join(post_training_files_path, image_file)
        img = cv2.imread(image_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        plt.figure(figsize=(10, 10), dpi=120)
        plt.imshow(img)
        plt.axis('off')
        plt.show()

# List of image files to display
image_files = [
    'confusion_matrix_normalized.png',
    'BoxF1 curve.png',
    'BoxP curve.png'
]
```

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** File, Edit, Selection, View, Go, Run, Terminal, Help, Search.
- Toolbar:** LAB_15 (1).ipynb, Generate, Code, Markdown, Run All, Clear All Outputs, Outline, Select Kernel.
- Section Header:** Validation Step.
- Code Cell:** Contains Python code for displaying images. The code imports os, cv2, and matplotlib.pyplot, defines a function display_images, and lists several image files to be displayed.
- Cell Number:** [16] at the bottom left.
- Bottom Status Bar:** Live Share, 0 14, Cell 1 of 20.

```
import os
import cv2
import matplotlib.pyplot as plt

def display_images(post_training_files_path, image_files):

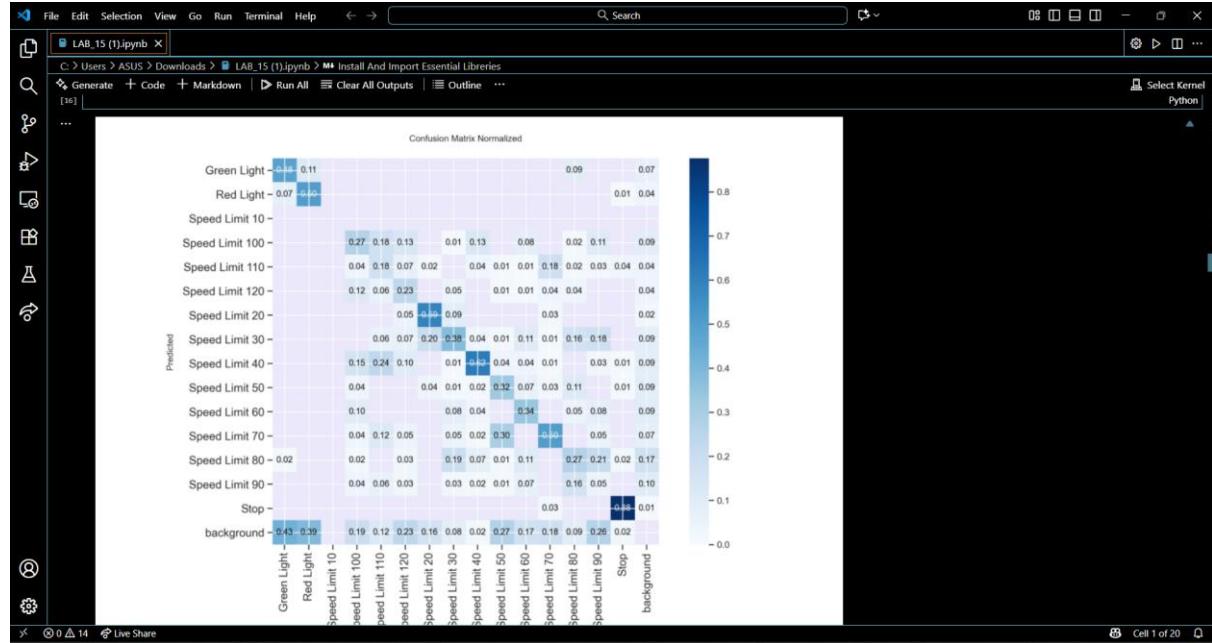
    for image_file in image_files:
        image_path = os.path.join(post_training_files_path, image_file)
        img = cv2.imread(image_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

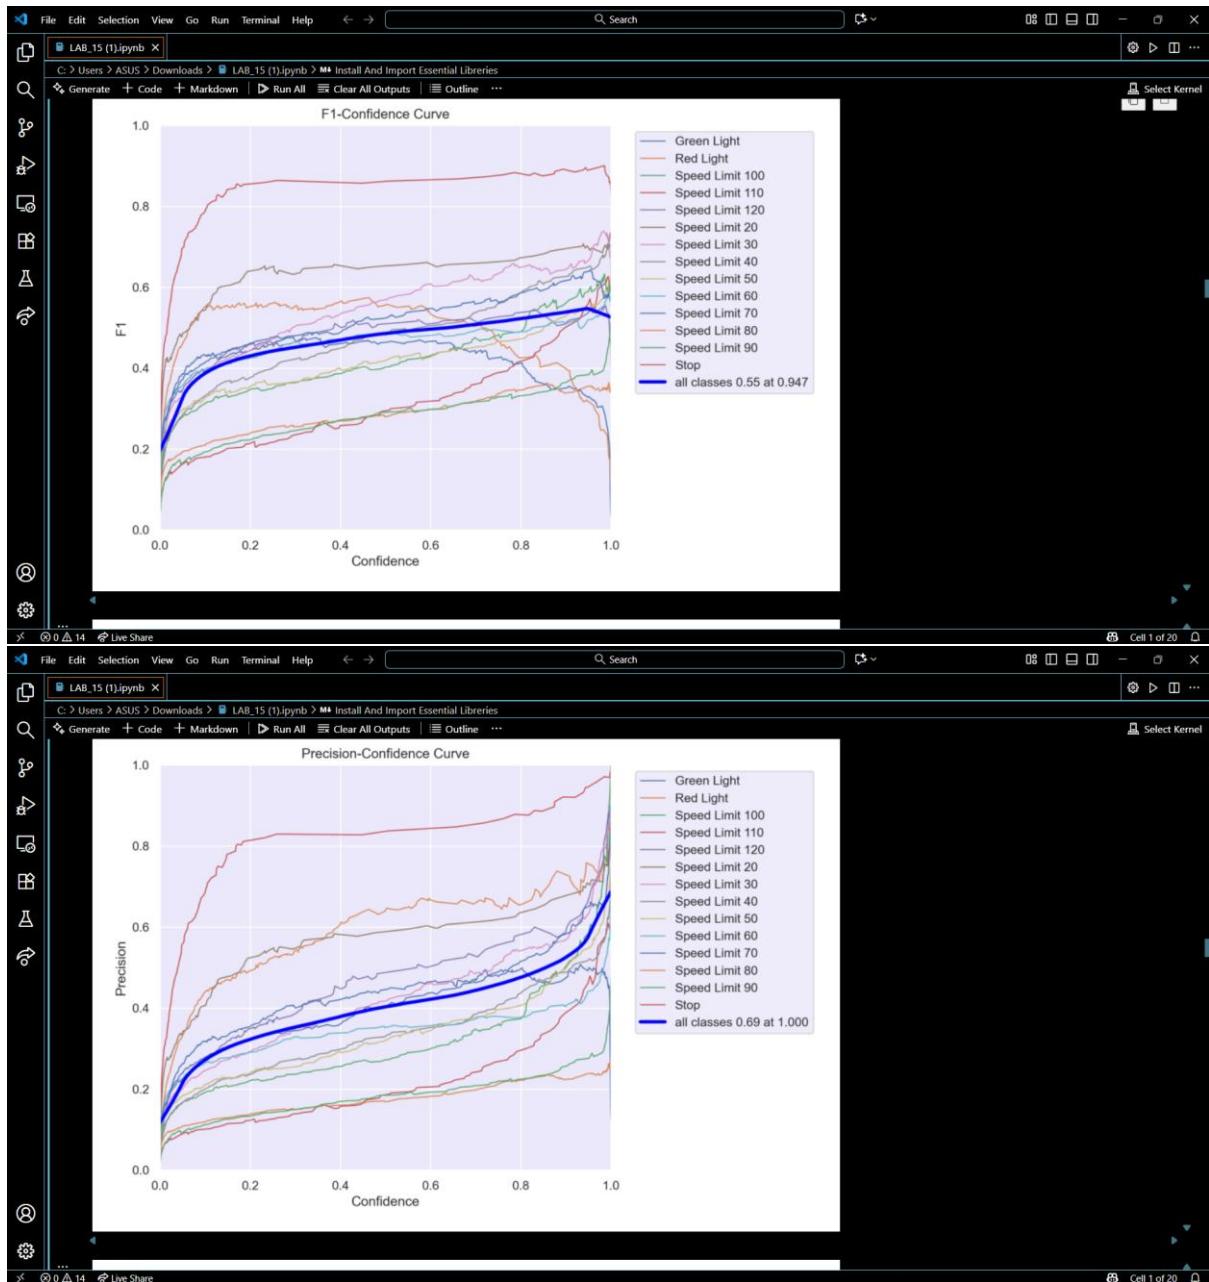
        plt.figure(figsize=(10, 10), dpi=120)
        plt.imshow(img)
        plt.axis('off')
        plt.show()

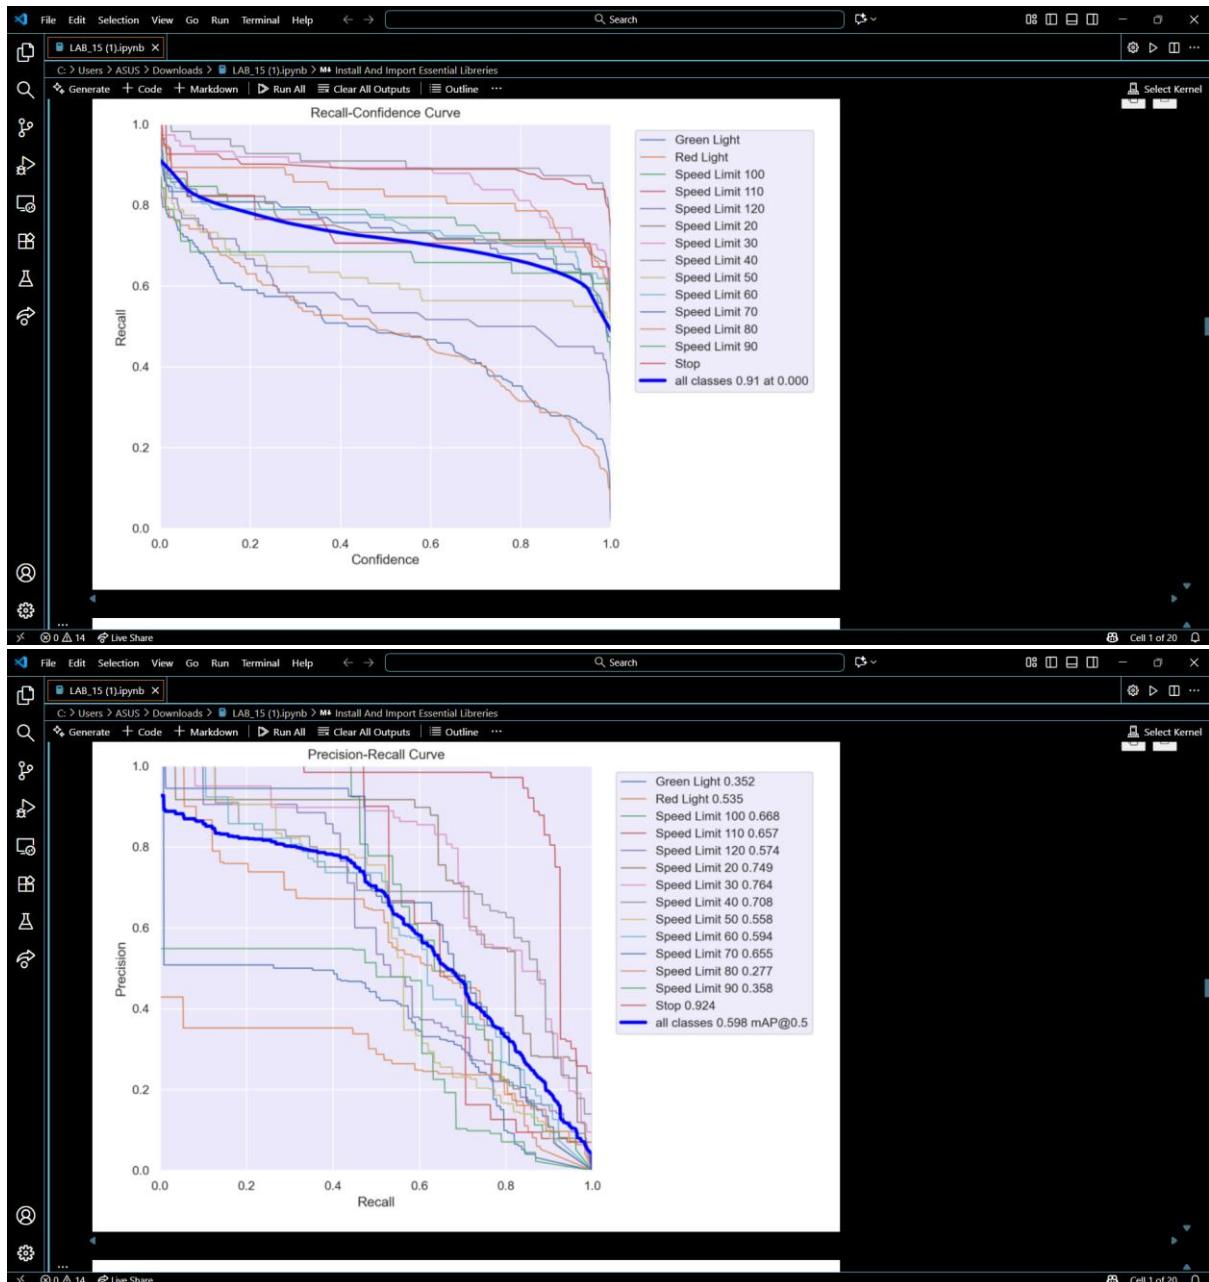
    # List of image files to display
    image_files = [
        'confusion_matrix_normalized.png',
        'BoxI_curve.png',
        'BoxP_curve.png',
        'BoxR_curve.png',
        'BoxR_curve.png',
        'results.png'
    ]

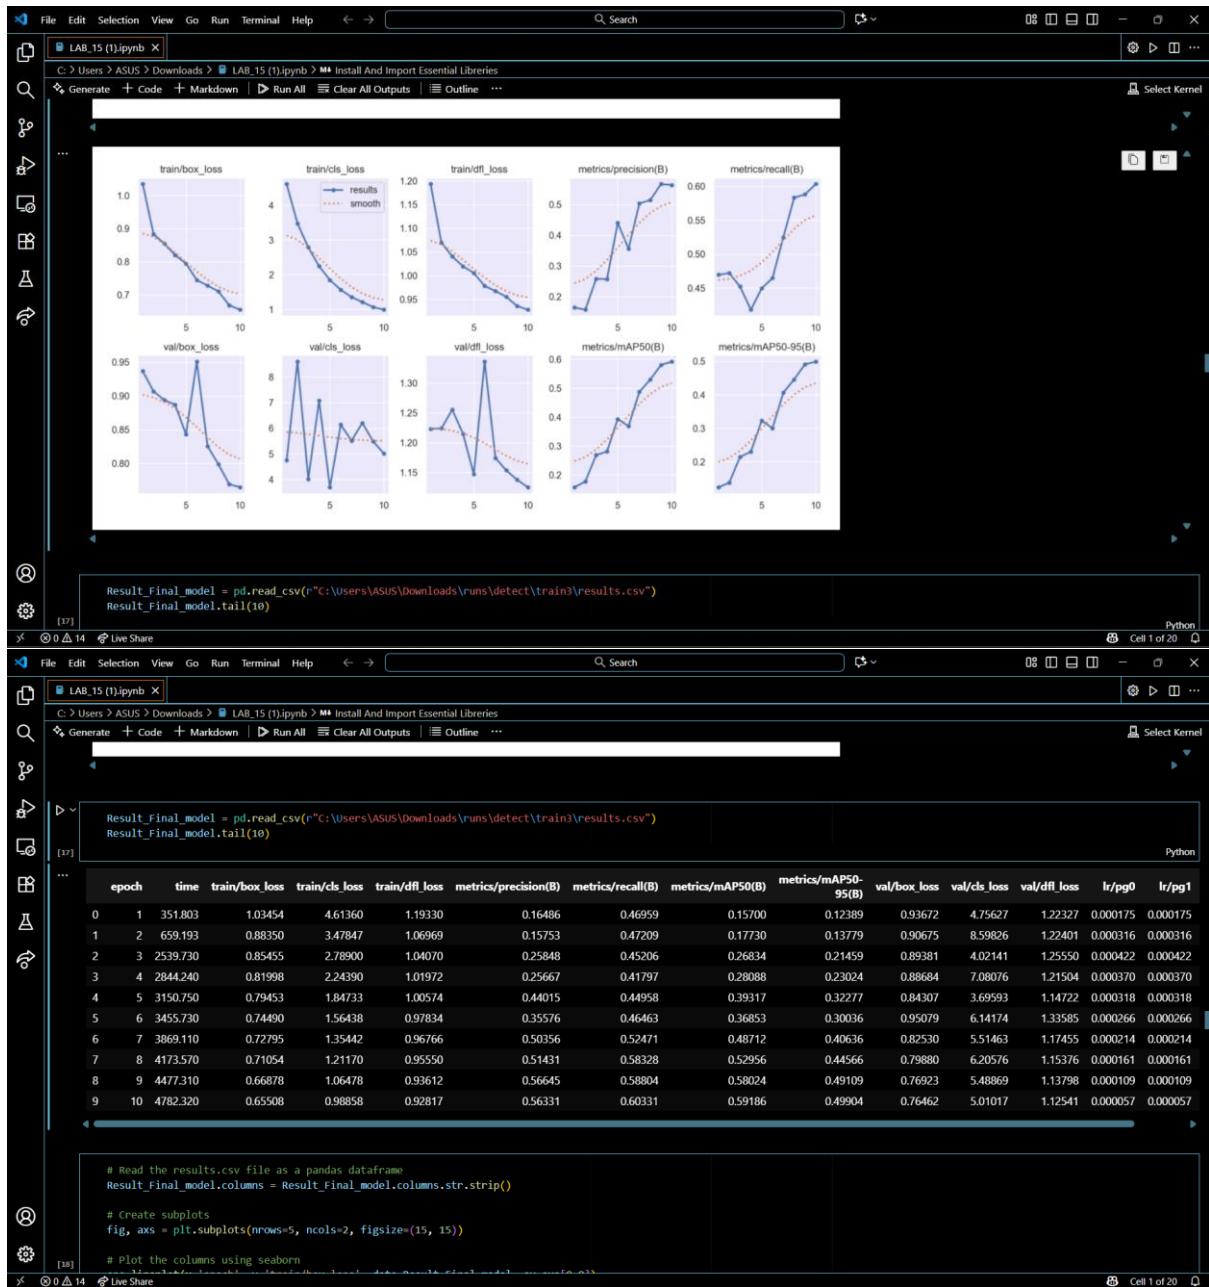
    # Path to the directory containing the images
    post_training_files_path = r"C:\Users\ASUS\Downloads\runs\detect\train3"

    # Display the images
    display_images(post_training_files_path, image_files)
```













File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search Select Kernel

C:\> Users > ASUS > Downloads > LAB_15 (1).ipynb > # Install And Import Essential Libraries

Generate + Code + Markdown | Run All Clear All Outputs | Outline ...

Validation of the Model By TestSet

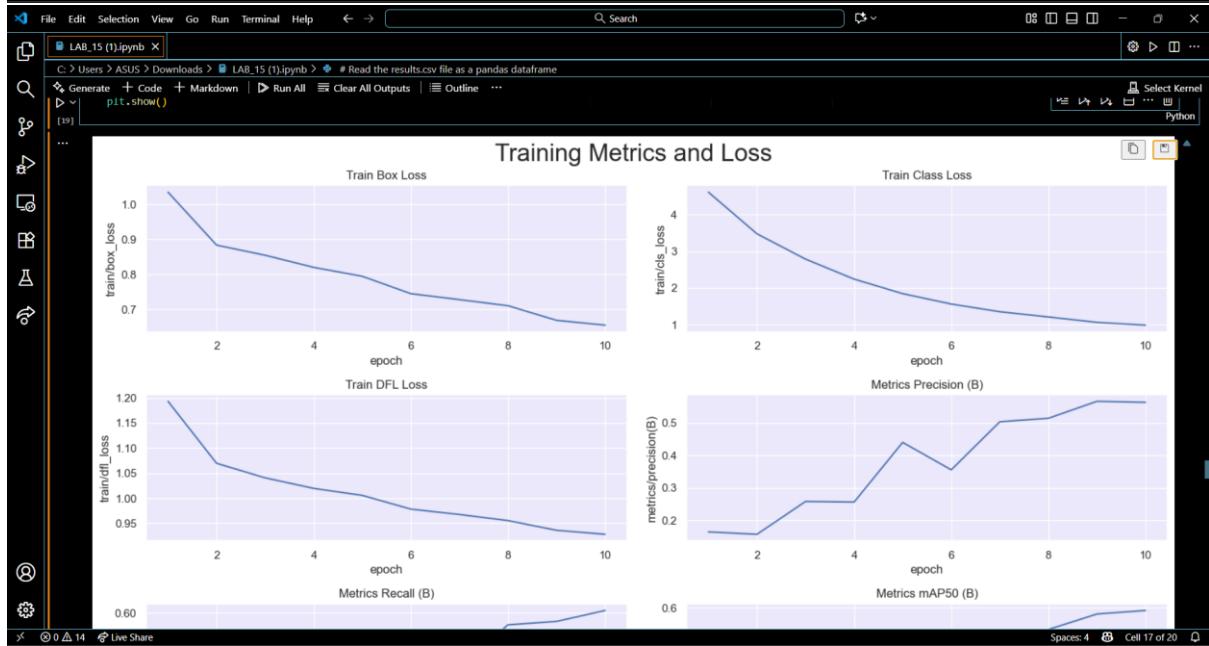
```
# Read the results.csv file as a pandas dataframe
Result_Final_model.columns = Result_Final_model.columns.str.strip()

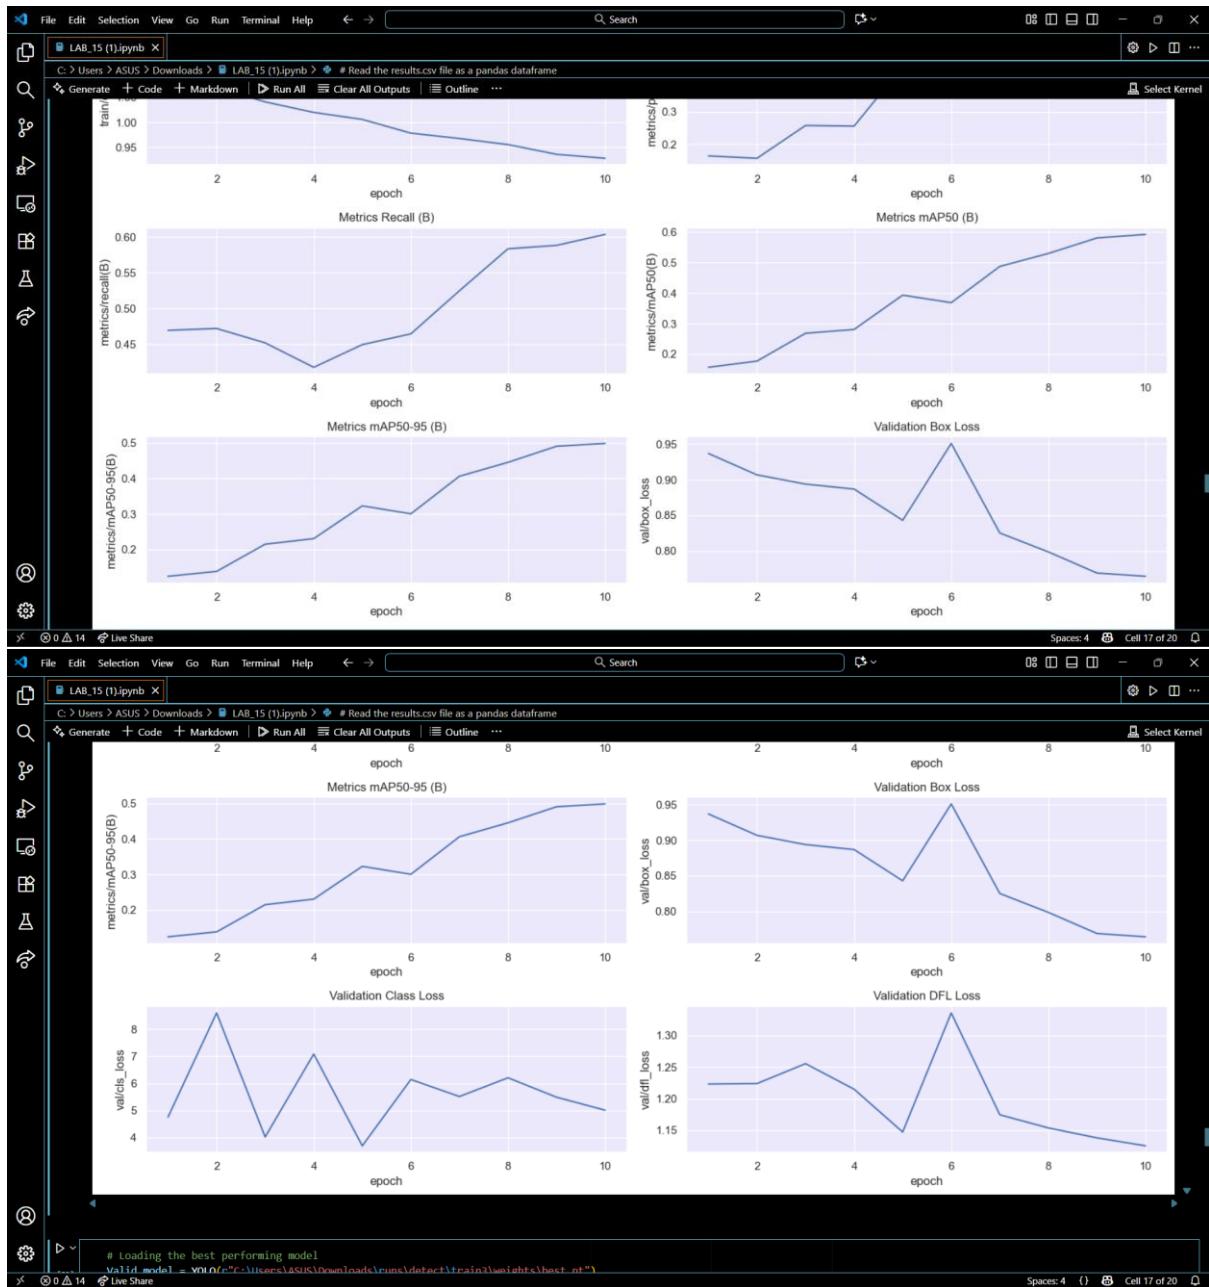
# Create subplots
fig, axs = plt.subplots(nrows=5, ncols=2, figsize=(15, 15))

# Plot the columns using seaborn
sns.lineplot(x="epoch", y="train/box_loss", data=Result_Final_model, ax=axs[0,0])
sns.lineplot(x="epoch", y="train/cls_loss", data=Result_Final_model, ax=axs[0,1])
sns.lineplot(x="epoch", y="train/dfl_loss", data=Result_Final_model, ax=axs[1,0])
sns.lineplot(x="epoch", y="metrics/precision(B)", data=Result_Final_model, ax=axs[1,1])
sns.lineplot(x="epoch", y="metrics/recall(B)", data=Result_Final_model, ax=axs[2,0])
sns.lineplot(x="epoch", y="metrics/mAP50(B)", data=Result_Final_model, ax=axs[2,1])
sns.lineplot(x="epoch", y="metrics/mAP50-95(B)", data=Result_Final_model, ax=axs[3,0])
sns.lineplot(x="epoch", y="val/box_loss", data=Result_Final_model, ax=axs[3,1])
sns.lineplot(x="epoch", y="val/cls_loss", data=Result_Final_model, ax=axs[4,0])
sns.lineplot(x="epoch", y="val/dfl_loss", data=Result_Final_model, ax=axs[4,1])

# Set titles and axis labels for each subplot
axs[0,0].set(title='Train Box loss')
axs[0,1].set(title='Train Class loss')
axs[1,0].set(title='Train DFL loss')
axs[1,1].set(title='Metrics Precision (B)')
axs[2,0].set(title='Metrics Recall (B)')
axs[2,1].set(title='Metrics mAP50 (B)')
axs[3,0].set(title='Validation Box loss')
axs[3,1].set(title='Validation Class loss')
axs[4,0].set(title='Validation DFL loss')
```

[18] Cell 1 of 20





LAB_15 (1).py

```

File Edit Selection View Go Run Terminal Help < > Search & Select Kernel
C:\Users\ASUS\Downloads> LAB_15 (1).py & # Read the results.csv file as a pandas dataframe
Generate + Code + Markdown | Run All Clear All Outputs | Outline ...
# Loading the best performing model
Valid_model = YOLO("C:\Users\ASUS\Downloads\runs\detect\train3\weights\best.pt")

# Evaluating the model on the validset
metrics = Valid_model.val(split = 'val')

# final results
print("precision(0): ", metrics.results_dict["metrics/precision(0)"])
print("metrics/recall(0): ", metrics.results_dict["metrics/recall(0)"])
print("metrics/mAP50(0): ", metrics.results_dict["metrics/mAP50(0)"])
print("metrics/mAP50-95(0): ", metrics.results_dict["metrics/mAP50-95(0)"])

...
Ultralytics 8.3.221 Python-3.10.11 torch-2.5.1+cu121 CUDA-0 (NVIDIA GeForce RTX 3050 Laptop GPU, 4096MiB)
Model summary (fused): 72 layers, 3,008,573 parameters, 0 gradients, 8.1 GFLOPs
Model Fast image access (ping: 0.16.0 ms, read: 31.912.8 MB/s, size: 11.2 KB)
val: Scanning D:\car\valid\labels\cache_... 801 images, 0 backgrounds, 0 corrupt: 100% 801/801 764.4Ki/s 0.0s
      Class   Images Instances Box(P) R mAP50 mAP50-95: 100% 51/51 5.8it/s 8.8s±0.1s
          all    801     944  0.565  0.694  0.625  0.529
        Green Light   87     122   0.51  0.247  0.351  0.192
        Red Light    74     108   0.758  0.203  0.536  0.318
      Speed Limit 100   52     52   0.581  0.596  0.668  0.614
      Speed Limit 110   17     17   0.457  0.706  0.668  0.615
      Speed Limit 120   60     60   0.663  0.45  0.577  0.49
      Speed Limit 20   56     56   0.709  0.695  0.791  0.69
      Speed Limit 30   71     74   0.622  0.703  0.766  0.687
      Speed Limit 40   53     55   0.516  0.873  0.716  0.619
      Speed Limit 50   68     71   0.547  0.549  0.58  0.486
      Speed Limit 60   76     76   0.433  0.632  0.635  0.556
      Speed Limit 70   78     78   0.649  0.628  0.678  0.589
      Speed Limit 80   56     56   0.246  0.696  0.332  0.272
      Speed Limit 90   38     38   0.275  0.632  0.516  0.449
      Stop         81     81   0.945  0.84  0.931  0.831

```

LAB_15 (1).py

```

File Edit Selection View Go Run Terminal Help < > Search & Select Kernel
C:\Users\ASUS\Downloads> LAB_15 (1).py & # Read the results.csv file as a pandas dataframe
Generate + Code + Markdown | Run All Clear All Outputs | Outline ...
precision(0): 0.569033325488223
metrics/recall(0): 0.6035039173819629
metrics/mAP50(0): 0.6246812693799372
metrics/mAP50-95(0): 0.5292130266848856

Making Predictions On Test Images

# Normalization function
def normalize_image(image):
|   return image / 255.0

# Image resizing function
def resize_image(image, size=(640, 640)):
|   return cv2.resize(image, size)

# Path to validation images
dataset_path = r'D:\car' # Place your dataset path here
valid_images_path = os.path.join(dataset_path, 'test', 'images')

# List of all jpg images in the directory
image_files = [file for file in os.listdir(valid_images_path) if file.endswith('.jpg')]

# Check if there are images in the directory
if len(image_files) > 0:
|   # Select 9 images at equal intervals
|   num_images = len(image_files)
|   step_size = max(1, num_images // 9) # Ensure the interval is at least 1
|   selected_images = [image_files[i] for i in range(0, num_images, step_size)]

```

LAB_15 (1).ipynb

```

C:\Users>ASUS>Downloads > LAB_15 (1).ipynb > # Read the results.csv file as a pandas dataframe
↳ Generate + Code + Markdown | ▶ Run All ━━ Clear All Outputs | ━━ Outline ... Select Kernel
Making Predictions On Test Images

# Normalization function
def normalize_image(image):
    return image / 255.0

# Image resizing function
def resize_image(image, size=(640, 640)):
    return cv2.resize(image, size)

# Path to validation images
dataset_path = r'D:\car' # Place your dataset path here
valid_images_path = os.path.join(dataset_path, 'test', 'images')

# List of all jpg images in the directory
image_files = [file for file in os.listdir(valid_images_path) if file.endswith('.jpg')]

# Check if there are images in the directory
if len(image_files) > 0:
    # Select 9 images at equal intervals
    num_images = len(image_files)
    step_size = max(1, num_images // 9) # Ensure the interval is at least 1
    selected_images = [image_files[i] for i in range(0, num_images, step_size)]

# Prepare subplots
fig, axes = plt.subplots(3, 3, figsize=(20, 21))
fig.suptitle('Validation Set Inferences', fontsize=24)

for i, ax in enumerate(axes.flatten()):
    if i < len(selected_images):
        image_path = os.path.join(valid_images_path, selected_images[i])

        # Load image
        image = cv2.imread(image_path)

        # Check if the image is loaded correctly
        if image is not None:
            # Resize image
            resized_image = resize_image(image, size=(640, 640))
            # Normalize image
            normalized_image = normalize_image(resized_image)

            # Convert the normalized image to uint8 data type
            normalized_image_uint8 = (normalized_image * 255).astype(np.uint8)

            # Predict with the model
            results = Valid_model.predict(source=normalized_image_uint8, imgsz=640, conf=0.5)

            # Plot image with labels
            annotated_image = results[0].plot(line_width=1)
            annotated_image_rgb = cv2.cvtColor(annotated_image, cv2.COLOR_BGR2RGB)
            ax.imshow(annotated_image_rgb)

        else:
            print(f"Failed to load image {image_path}")
            ax.axis('off')

    plt.tight_layout()

```

Spaces: 4 Cell 17 of 20

LAB_15 (1).ipynb

```

C:\Users>ASUS>Downloads > LAB_15 (1).ipynb > # Read the results.csv file as a pandas dataframe
↳ Generate + Code + Markdown | ▶ Run All ━━ Clear All Outputs | ━━ Outline ... Select Kernel
selected_images = [image_files[i] for i in range(0, num_images, step_size)]

# Prepare subplots
fig, axes = plt.subplots(3, 3, figsize=(20, 21))
fig.suptitle('Validation Set Inferences', fontsize=24)

for i, ax in enumerate(axes.flatten()):
    if i < len(selected_images):
        image_path = os.path.join(valid_images_path, selected_images[i])

        # Load image
        image = cv2.imread(image_path)

        # Check if the image is loaded correctly
        if image is not None:
            # Resize image
            resized_image = resize_image(image, size=(640, 640))
            # Normalize image
            normalized_image = normalize_image(resized_image)

            # Convert the normalized image to uint8 data type
            normalized_image_uint8 = (normalized_image * 255).astype(np.uint8)

            # Predict with the model
            results = Valid_model.predict(source=normalized_image_uint8, imgsz=640, conf=0.5)

            # Plot image with labels
            annotated_image = results[0].plot(line_width=1)
            annotated_image_rgb = cv2.cvtColor(annotated_image, cv2.COLOR_BGR2RGB)
            ax.imshow(annotated_image_rgb)

        else:
            print(f"Failed to load image {image_path}")
            ax.axis('off')

    plt.tight_layout()

```

Spaces: 4 Cell 17 of 20

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search Select Kernel Python

```

LAB_15 (1).ipynb x
C:\Users\ASUS\Downloads> LAB_15 (1).ipynb > # Read the results.csv file as a pandas dataframe
... Generate + Code + Markdown | Run All Clear All Outputs Outline ...
| ax.axis('off')
| plt.tight_layout()
| plt.show()

[22] ...
0: 640x640 1 Speed Limit 30, 1 Speed Limit 60, 58.5ms
Speed: 5.0ms preprocess, 58.5ms inference, 3.6ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Speed Limit 30, 58.1ms
Speed: 4.2ms preprocess, 58.1ms inference, 7.2ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Speed Limit 70, 8.1ms
Speed: 1.3ms preprocess, 8.1ms inference, 1.3ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Stop, 8.1ms
Speed: 1.6ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Speed Limit 70s, 8.3ms
Speed: 1.1ms preprocess, 8.3ms inference, 1.3ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Green Lights, 9.8ms
Speed: 1.3ms preprocess, 9.8ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Red Light, 8.1ms
Speed: 1.1ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 (no detections), 8.1ms
Speed: 1.3ms preprocess, 8.1ms inference, 0.7ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Speed Limit 60, 1 Speed Limit 80, 1 Speed Limit 90, 8.0ms
Speed: 1.2ms preprocess, 8.0ms inference, 1.6ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Stop, 8.1ms
Speed: 1.6ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Speed Limit 70s, 8.3ms
Speed: 1.1ms preprocess, 8.3ms inference, 1.3ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Green Lights, 9.8ms
Speed: 1.3ms preprocess, 9.8ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Red Light, 8.1ms
Speed: 1.1ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 (no detections), 8.1ms
Speed: 1.3ms preprocess, 8.1ms inference, 0.7ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Speed Limit 60, 1 Speed Limit 80, 1 Speed Limit 90, 8.0ms
Speed: 1.2ms preprocess, 8.0ms inference, 1.6ms postprocess per image at shape (1, 3, 640, 640)

```

Live Share

File Edit Selection View Go Run Terminal Help ⏪ ⏩ Search Select Kernel Python

```

LAB_15 (1).ipynb x
C:\Users\ASUS\Downloads> LAB_15 (1).ipynb > # Read the results.csv file as a pandas dataframe
... Generate + Code + Markdown | Run All Clear All Outputs Outline ...
0: 640x640 1 Stop, 8.1ms
Speed: 1.6ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Speed Limit 70s, 8.3ms
Speed: 1.1ms preprocess, 8.3ms inference, 1.3ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Green Lights, 9.8ms
Speed: 1.3ms preprocess, 9.8ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Red Light, 8.1ms
Speed: 1.1ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 (no detections), 8.1ms
Speed: 1.3ms preprocess, 8.1ms inference, 0.7ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Speed Limit 60, 1 Speed Limit 80, 1 Speed Limit 90, 8.0ms
Speed: 1.2ms preprocess, 8.0ms inference, 1.6ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Stop, 8.1ms
Speed: 1.6ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Speed Limit 70s, 8.3ms
Speed: 1.1ms preprocess, 8.3ms inference, 1.3ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 2 Green Lights, 9.8ms
Speed: 1.3ms preprocess, 9.8ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Red Light, 8.1ms
Speed: 1.1ms preprocess, 8.1ms inference, 1.4ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 (no detections), 8.1ms
Speed: 1.3ms preprocess, 8.1ms inference, 0.7ms postprocess per image at shape (1, 3, 640, 640)

0: 640x640 1 Speed Limit 60, 1 Speed Limit 80, 1 Speed Limit 90, 8.0ms
Speed: 1.2ms preprocess, 8.0ms inference, 1.6ms postprocess per image at shape (1, 3, 640, 640)

```

Validation Set Inferences

